

# HOMEWORK 6: TEXT CLASSIFICATION

In this homework, you will create models to classify texts from TRUE call-center. There are two classification tasks:

1. Action Classification: Identify which action the customer would like to take (e.g. enquire, report, cancel)
2. Object Classification: Identify which object the customer is referring to (e.g. payment, true money, internet, roaming)

We will focus only on the Object Classification task for this homework.

In this homework, you are asked to compare different text classification models in terms of accuracy and inference time.

You will need to build 3 different models.

1. A model based on tf-idf
2. A model based on MUSE
3. A model based on wangchanBERTa

**You will be asked to submit 3 different files (.pdf from .ipynb) that do the 3 different models. Finally, answer the accuracy and runtime numbers in MCV.**

This homework is quite free form, and your answer may vary. We hope that the processing during the course of this assignment will make you think more about the design choices in text classification.

```
In [46]: # !wget --no-check-certificate https://www.dropbox.com/s/37u83g55p19kvrL/clean-p
```

```
In [47]: # !pip install pythainlp
```

## Import Libs

```
In [48]: # %matplotlib inline
# import pandas
# import sklearn
# import numpy as np
# import matplotlib.pyplot as plt
# import pandas as pd

# from torch.utils.data import Dataset
# from IPython.display import display
# from collections import defaultdict
# from sklearn.metrics import accuracy_score
# from sklearn.model_selection import train_test_split
# from pythainlp.tokenize import word_tokenize
# from sklearn.feature_extraction.text import TfidfVectorizer
# from sklearn.linear_model import LogisticRegression
# from sklearn.pipeline import Pipeline
```

```
# from pythainlp.corpus.common import thai_stopwords
# import time
# import torch
# from transformers import AutoTokenizer, AutoModelForSequenceClassification, Tr
# from datasets import Dataset
# from sklearn.preprocessing import LabelEncoder
```

## Loading data

First, we load the data from disk into a Dataframe.

A Dataframe is essentially a table, or 2D-array/Matrix with a name for each column.

```
In [49]: data_df = pd.read_csv('clean-phone-data-for-students.csv')
```

Let's preview the data.

```
In [50]: # Show the top 5 rows
display(data_df.head())
# Summarize the data
data_df.describe()
```

|   | Sentence Utterance                               | Action  | Object       |
|---|--|---------|--------------|
| 0 | <PHONE_NUMBER_REMOVED> ผมไปจ่ายเงินที่ Counte... | enquire | payment      |
| 1 | internet ยังความเร็วอยู่เท่าไรครับ               | enquire | package      |
| 2 | ตะกี้ไปชำระค่าบริการไปแล้ว แต่ยังใช้งานไม่ได้... | report  | suspend      |
| 3 | พี่คะยังใช้ internet ไม่ได้เลยคะ เป็นเครือ...    | enquire | internet     |
| 4 | ฮาโหล คะ พอดีว่าเมื่อวานเปิดซิมทรูมูฟ แต่มันโ... | report  | phone_issues |

```
Out[50]:
```

|        | Sentence Utterance | Action  | Object  |
|--------|--------------------|---------|---------|
| count  | 16175              | 16175   | 16175   |
| unique | 13389              | 10      | 33      |
| top    | บริการอื่นๆ        | enquire | service |
| freq   | 97                 | 10377   | 2525    |

## Data cleaning

We call the DataFrame.describe() again. Notice that there are 33 unique labels/classes for object and 10 unique labels for action that the model will try to predict. But there are unwanted duplications e.g. ldd,idd,lotalty\_card,Lotalty\_card

Also note that, there are 13389 unique sentence utterances from 16175 utterances. You have to clean that too!

## #TODO 0.1:

You will have to remove unwanted label duplications as well as duplications in text inputs. Also, you will have to trim out unwanted whitespaces from the text inputs. This shouldn't be too hard, as you have already seen it in the demo.

```
In [51]: display(data_df.describe())
display(data_df.Object.unique())
display(data_df.Action.unique())
```

|        | Sentence Utterance | Action  | Object  |
|--------|--------------------|---------|---------|
| count  | 16175              | 16175   | 16175   |
| unique | 13389              | 10      | 33      |
| top    | บริการอื่นๆ        | enquire | service |
| freq   | 97                 | 10377   | 2525    |

```
array(['payment', 'package', 'suspend', 'internet', 'phone_issues',
       'service', 'nonTrueMove', 'balance', 'detail', 'bill', 'credit',
       'promotion', 'mobile_setting', 'iservice', 'roaming', 'truemoney',
       'information', 'lost_stolen', 'balance_minutes', 'idd',
       'TrueMoney', 'garbage', 'Payment', 'IDD', 'ringtone', 'Idd',
       'rate', 'loyalty_card', 'contact', 'officer', 'Balance', 'Service',
       'Loyalty_card'], dtype=object)
array(['enquire', 'report', 'cancel', 'Enquire', 'buy', 'activate',
       'request', 'Report', 'garbage', 'change'], dtype=object)
```

```
In [52]: data_df.columns
```

```
Out[52]: Index(['Sentence Utterance', 'Action', 'Object'], dtype='object')
```

```
In [53]: cols = ["Sentence Utterance", "Object"]
data_df = data_df[cols]
data_df.columns = ["input", "raw_label"]

data_df["clean_label"] = data_df["raw_label"].str.lower().copy()
data_df.drop("raw_label", axis=1, inplace=True)

data_df["input"] = data_df["input"].str.strip()

data_df = data_df.drop_duplicates(subset=["input"], keep='first')
```

```
In [54]: display(data_df["clean_label"].unique())
display(data_df.describe())
display(data_df.head())
```

```
array(['payment', 'package', 'suspend', 'internet', 'phone_issues',
       'service', 'nontruemove', 'balance', 'detail', 'bill', 'credit',
       'promotion', 'mobile_setting', 'iservice', 'roaming', 'truemoney',
       'information', 'lost_stolen', 'balance_minutes', 'idd', 'garbage',
       'ringtone', 'rate', 'loyalty_card', 'contact', 'officer'],
      dtype=object)
```

|        | input                                | clean_label |
|--------|--------------------------------------|-------------|
| count  | 13367                                | 13367       |
| unique | 13367                                | 26          |
| top    | สอบถามโปรโมชั่นปัจจุบันที่ใช้อยู่ค่ะ |             |
| freq   | 1                                    | 2108        |

|   | input   | clean_label  |
|---|---|--------------|
| 0 | <PHONE_NUMBER_REMOVED> ผมไปจ่ายเงินที่ Counter...   | payment      |
| 1 | internet ยังความเร็วอยู่เท่าไรครับ                  | package      |
| 2 | ตะกี้ไปชำระค่าบริการไปแล้ว แต่ยังไม่ทำงานไม่ได้ ค่ะ | suspend      |
| 3 | พี่คะยังใช้ internet ไม่ได้เลยคะ เป็นเครื่อง...     | internet     |
| 4 | ฮัลโหล ค่ะ พอดีว่าเมื่อวานเปิดซิมทรูฟ แต่มันโท...   | phone_issues |

Split data into train, validation, and test sets (normally the ratio will be 80:10:10 , respectively). We recommend to use `train_test_split` from `scikit-learn` to split the data into train, validation, test set.

In addition, it should split the data that distribution of the labels in train, validation, test set are similar. There is **stratify** option to handle this issue.

[https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.train\\_test\\_split.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html)

Make sure the same data splitting is used for all models.

```
In [55]: data_x = np.array(list(data_df["input"]))
data_y_tmp = np.array(list(data_df["clean_label"]))
data_y = []

map_label_num = {y.strip():i for i,y in enumerate(list(data_df["clean_label"].unique()))}
map_num_label = {i:y.strip() for i,y in enumerate(list(data_df["clean_label"].unique()))}

for i in range(len(data_y_tmp)):
    data_y.append(int(map_label_num[data_y_tmp[i]]))
data_y = np.array(data_y)
print(len(data_y))
```

13367

```
In [56]: unique, counts = np.unique(data_y, return_counts=True)
valid_classes = unique[counts >= 10]
valid_indices = np.isin(data_y, valid_classes)
data_x, data_y = data_x[valid_indices], data_y[valid_indices]
```

```
In [57]: X_train, X_temp, y_train, y_temp = train_test_split(data_x, data_y, test_size=0.1)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.50,
print("Train size:", len(X_train))
```

```
print("Validation size:", len(X_val))  
print("Test size:", len(X_test))
```

Train size: 10690

Validation size: 1336

Test size: 1337

## Model 3 WangchanBERTa

We ask you to train a WangchanBERTa-based model.

We recommend you use the thaixtransformers fork (which we used in the PoS homework). <https://github.com/PyThaiNLP/thaixtransformers>

The structure of the code will be very similar to the PoS homework. You will also find the huggingface [tutorial](#) useful. Or you can also add a softmax layer by yourself just like in the previous homework.

Which WangchanBERTa model will you use? Why? (Don't forget to clean your text accordingly).

**Ans:**

In [58]: `!pip install wandb`

Requirement already satisfied: wandb in /usr/local/lib/python3.10/dist-packages (0.19.1)

Requirement already satisfied: click!=8.0.0,>=7.1 in /usr/local/lib/python3.10/dist-packages (from wandb) (8.1.7)

Requirement already satisfied: docker-pycreds>=0.4.0 in /usr/local/lib/python3.10/dist-packages (from wandb) (0.4.0)

Requirement already satisfied: gitpython!=3.1.29,>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from wandb) (3.1.43)

Requirement already satisfied: platformdirs in /usr/local/lib/python3.10/dist-packages (from wandb) (4.3.6)

Requirement already satisfied: protobuf!=4.21.0,!5.28.0,<6,>=3.19.0 in /usr/local/lib/python3.10/dist-packages (from wandb) (3.20.3)

Requirement already satisfied: psutil>=5.0.0 in /usr/local/lib/python3.10/dist-packages (from wandb) (5.9.5)

Requirement already satisfied: pydantic<3,>=2.6 in /usr/local/lib/python3.10/dist-packages (from wandb) (2.11.0a1)

Requirement already satisfied: pyyaml in /usr/local/lib/python3.10/dist-packages (from wandb) (6.0.2)

Requirement already satisfied: requests<3,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from wandb) (2.32.3)

Requirement already satisfied: sentry-sdk>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from wandb) (2.19.2)

Requirement already satisfied: setproctitle in /usr/local/lib/python3.10/dist-packages (from wandb) (1.3.4)

Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from wandb) (75.1.0)

Requirement already satisfied: typing-extensions<5,>=4.4 in /usr/local/lib/python3.10/dist-packages (from wandb) (4.12.2)

Requirement already satisfied: six>=1.4.0 in /usr/local/lib/python3.10/dist-packages (from docker-pycreds>=0.4.0->wandb) (1.17.0)

Requirement already satisfied: gitdb<5,>=4.0.1 in /usr/local/lib/python3.10/dist-packages (from gitpython!=3.1.29,>=1.0.0->wandb) (4.0.11)

Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.10/dist-packages (from pydantic<3,>=2.6->wandb) (0.7.0)

Requirement already satisfied: pydantic-core==2.28.0 in /usr/local/lib/python3.10/dist-packages (from pydantic<3,>=2.6->wandb) (2.28.0)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.0.0->wandb) (3.4.1)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.0.0->wandb) (3.10)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.0.0->wandb) (2.3.0)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.0.0->wandb) (2025.1.31)

Requirement already satisfied: smmap<6,>=3.0.1 in /usr/local/lib/python3.10/dist-packages (from gitdb<5,>=4.0.1->gitpython!=3.1.29,>=1.0.0->wandb) (5.0.1)

```
In [59]: from kaggle_secrets import UserSecretsClient
import wandb
user_secrets = UserSecretsClient()

my_secret = user_secrets.get_secret("wandb_api_key")

wandb.login(key=my_secret)
```

wandb: WARNING Calling wandb.login() after wandb.init() has no effect.

Out[59]: True

```

In [60]: import torch
import pandas as pd
from transformers import AutoTokenizer, AutoModelForSequenceClassification, TrainingArguments
from datasets import Dataset
from sklearn.preprocessing import LabelEncoder

# Load tokenizer and model
model_name = "airesearch/wangchanberta-base-att-spm-uncased"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForSequenceClassification.from_pretrained(model_name, num_labels=10)

# Encode labels
label_encoder = LabelEncoder()
y_train_enc = label_encoder.fit_transform(y_train)
y_val_enc = label_encoder.transform(y_val)
y_test_enc = label_encoder.transform(y_test)

# Tokenize data
def tokenize_function(examples):
    return tokenizer(examples["text"], padding="max_length", truncation=True, max_length=512)

train_data = Dataset.from_dict({"text": X_train, "label": y_train_enc}).map(tokenize_function)
val_data = Dataset.from_dict({"text": X_val, "label": y_val_enc}).map(tokenize_function)
test_data = Dataset.from_dict({"text": X_test, "label": y_test_enc}).map(tokenize_function)

# Define training arguments
training_args = TrainingArguments(
    output_dir="./results", # Keep output directory for saving checkpoints
    run_name="wangchanberta_classification", # Set a different name for W&B
    eval_strategy="epoch",
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    num_train_epochs=3,
    save_strategy="epoch",
    save_total_limit=1,
    logging_dir="./logs",
    logging_steps=50,
    load_best_model_at_end=True
)

# Trainer
def compute_metrics(eval_pred):
    logits, labels = eval_pred
    predictions = np.argmax(logits, axis=-1) # Get the predicted class
    acc = accuracy_score(labels, predictions) # Compute accuracy
    return {"accuracy": acc}

trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_data,
    eval_dataset=val_data,
    compute_metrics=compute_metrics # Add the metrics function here
)

# Train model
start = time.time()

```

```
trainer.train()
end = time.time()
```

Some weights of CamembertForSequenceClassification were not initialized from the model checkpoint at airesearch/wangchanberta-base-att-spm-uncased and are newly initialized: ['classifier.dense.bias', 'classifier.dense.weight', 'classifier.out\_proj.bias', 'classifier.out\_proj.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
Map: 0%|          | 0/10690 [00:00<?, ? examples/s]
Map: 0%|          | 0/1336 [00:00<?, ? examples/s]
Map: 0%|          | 0/1337 [00:00<?, ? examples/s]
```

[2007/2007 07:05, Epoch 3/3]

| Epoch | Training Loss | Validation Loss | Accuracy |
|-------|---------------|-----------------|----------|
| 1     | 1.338700      | 1.297383        | 0.618263 |
| 2     | 1.149300      | 1.343262        | 0.589072 |
| 3     | 0.610100      | 0.750168        | 0.773204 |

```
In [66]: train_results = trainer.evaluate(train_data)
val_results = trainer.evaluate(val_data)
test_results = trainer.evaluate(test_data)
```

```
In [67]: print(f"Training Time: {end - start:.4f} seconds")
print(f"Train Accuracy: {train_results['eval_accuracy']:.4f}")
print(f"Validation Accuracy: {val_results['eval_accuracy']:.4f}")
print(f"Test Accuracy: {test_results['eval_accuracy']:.4f}")
```

```
Training Time: 426.8042 seconds
Train Accuracy: 0.8536
Validation Accuracy: 0.7732
Test Accuracy: 0.7644
```

## Comparison

After you have completed the 3 models, compare the accuracy, ease of implementation, and inference speed (from cleaning, tokenization, till model compute) between the three models in mycourseville.

### Model1

- Training time: 3.0621 seconds
- Train Accuracy: 0.7650
- Validation Accuracy: 0.6939
- Test Accuracy: 0.6971

### Model 2

- Training Time: 2.2055 seconds
- Train Accuracy: 0.7351
- Validation Accuracy: 0.7118



- Test Accuracy: 0.7023

## Model 3

- Training Time: 426.8042 seconds
- Train Accuracy: 0.8536
- Validation Accuracy: 0.7732
- Test Accuracy: 0.7644

## ANS

WangchanBERTa ดีที่สุด เพราะ มีaccuracyสูงสุดและเรากำลังทำCallCenterChatbotซึ่งไม่จำเป็นต้องเร็วมากนัก